

# Topic modeling demo

## Load necessary libraries

```
library(tidyverse)
library(tidymodels)
library(tidylo)
library(stm)
library(purrr)
library(RColorBrewer)
library(ggwordcloud)
library(ggrepel)
library(patchwork)

# devtools::install_github("juliasilge/tidytext", force = TRUE)
library(tidytext)

theme_set(theme_bw())
```

## Load data: publicly available text datasets for the Office dialogue across all seasons:

```
library(schrute)
df_office <- schrute::theoffice
speaking_characters <- df_office %>%
  count(character, sort = TRUE) %>%

  # select only characters with at least 500 lines across all seasons
  filter(n > 500) %>%
  pull(character)
```

## Unnest tokens: restructure as one-token-per-row (each word is a separate row in a tibble)

```
df_tokens <- df_office %>%

  # filter for our speaking characters
  filter(character %in% speaking_characters) %>%

  group_by(character) %>%
  unnest_tokens(word, text) %>%

  # remove stop words (uninteresting words like of, from, and, the, etc.)
  anti_join(get_stopwords(), by = join_by(word))
```

```
head(df_tokens)
```

```
## # A tibble: 6 x 2
## # Groups:   character [1]
##   character word
##   <chr>      <chr>
## 1 Michael   right
## 2 Michael   jim
## 3 Michael   quarterlies
## 4 Michael   look
## 5 Michael   good
## 6 Michael   things
```

Create a sparse matrix: each row is a corpus (book) and each column is a word

```
sparse_mat <- df_tokens %>%
  # count number of times a word is used by each character
  dplyr::count(character, word) %>%
  # filter for words that appear at least 3 times
  filter(n > 3) %>%
  # create sparse matrix
  cast_sparse(character, word, n)

dim(sparse_mat)
```

```
## [1] 19 3016
```

```
sparse_mat[1:10, 1:10]
```

```
## 10 x 10 sparse Matrix of class "dgCMatrix"
##
## Andy    4 5 4 6 5 8 6 4 4 50
## Angela . . . . . . . . . 7
## Darryl . . . . . . . . . 8
## Dwight . . . 5 12 8 . 6 8 24
## Erin    . . . . . . . . . 5
## Holly   . . . . . . . . . 4
## Jan     . . . . . . . . . 7
## Jim     . 6 7 6 . 13 . . . 111
## Kelly   . . 4 . . . . . . 18
## Kevin   . . . . . . . . . 7
```

Fit topic model using Structural Topic Models (<http://www.structuraltopicmodel.com/>)

```
set.seed(0306)
topic_model <- stm(sparse_mat,
  K = 4, # K is the number of 'topics' (clusters)
  verbose = FALSE)

summary(topic_model)
```

## A topic model with 4 topics, 19 documents and a 3016 word dictionary.

```
## Topic 1 Top Words:
##   Highest Prob: just, oh, know, yeah, like, michael, get
##   FREX: michael, andy, um, mean, sorry, hi, kevin
##   Lift: channel, foster, cuz, executive, ruff, bloody, blake
##   Score: ethics, senator, cece, mural, pens, pum, confused
## Topic 2 Top Words:
##   Highest Prob: just, know, oh, yeah, right, like, hey
##   FREX: tuna, whoa, na, mike, gotta, pretty, nope
##   Lift: agent, america's, andrew, anger, audition, ba, banjo
##   Score: tuna, nard, jakey, wimowheh, philly, na, andrew
## Topic 3 Top Words:
##   Highest Prob: oh, know, just, okay, like, get, jim
##   FREX: schrute, mose, hay, farm, sensei, beet, male
##   Lift: hay, immediately, knife, wake, weak, 107th, 6
##   Score: spell, hay, mose, sensei, male, belsnickel, battle
## Topic 4 Top Words:
##   Highest Prob: know, just, oh, okay, like, well, right
##   FREX: beep, award, trouble, somebody, everybody, 300, rabies
##   Lift: 200, 4, 500, 645, a.j, ability, accountant
##   Score: packer, anxiety, beep, award, boom, trouble, anybody
```

**lift** = frequency divided by frequency in other topics **FREX** weights words by frequency and exclusivity to the topic

Ta-da!!

But... how can we determine the optimal number of topics for our dataset?

```
df_many_models <- tibble(K = seq(3, 10, by = 1)) %>%
  mutate(topic_model = map (
    .x = K,
    .f = ~ stm(sparse_mat,
      K = .x,
      verbose = FALSE)
  ))

head(df_many_models)
```

```
## # A tibble: 6 x 2
##   K topic_model
```

```
##    <dbl> <list>
## 1      3 <STM>
## 2      4 <STM>
## 3      5 <STM>
## 4      6 <STM>
## 5      7 <STM>
## 6      8 <STM>
```

Now that we've fit all these topic models with different numbers of topics, we can explore how many topics are appropriate/good/"best"

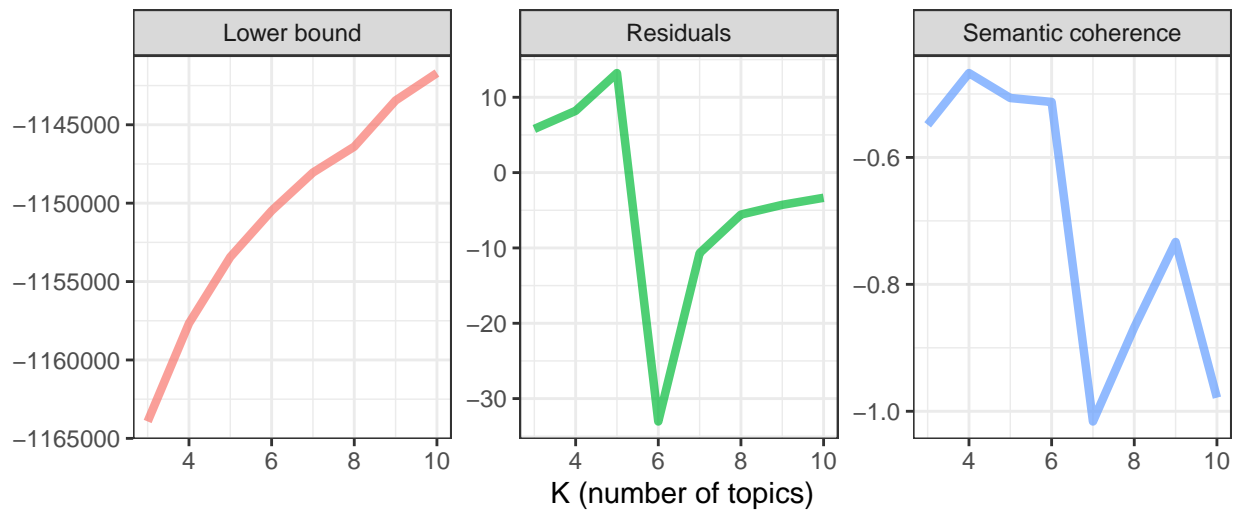
```
df_diagnostics <- df_many_models %>%
  mutate(exclusivity = map(topic_model, exclusivity),
         semantic_coherence = map(topic_model, semanticCoherence, sparse_mat),
         residual = map(topic_model, checkResiduals, sparse_mat),
         bound = map_dbl(topic_model, function(x) max(x$convergence$bound)),
         lfact = map_dbl(topic_model, function(x) lfactorial(x$settings$dim$K)),
         lbound = bound + lfact,
         iterations = map_dbl(topic_model, function(x) length(x$convergence$bound)))

head(df_diagnostics)
```

```
## # A tibble: 6 x 9
##       K topic_model exclusivity semantic_coherence residual      bound lfact
##   <dbl> <list>      <list>      <list>          <list>      <dbl> <dbl>
## 1     3 <STM>      <dbl [3]>    <dbl [3]>      <named list> -1163925.  1.79
## 2     4 <STM>      <dbl [4]>    <dbl [4]>      <named list> -1157667.  3.18
## 3     5 <STM>      <dbl [5]>    <dbl [5]>      <named list> -1153438.  4.79
## 4     6 <STM>      <dbl [6]>    <dbl [6]>      <named list> -1150492.  6.58
## 5     7 <STM>      <dbl [7]>    <dbl [7]>      <named list> -1148057.  8.53
## 6     8 <STM>      <dbl [8]>    <dbl [8]>      <named list> -1146407. 10.6
## # i 2 more variables: lbound <dbl>, iterations <dbl>
```

```
df_diagnostics %>%
  transmute(K,
            `Lower bound` = lbound,
            Residuals = map_dbl(residual, "dispersion"),
            `Semantic coherence` = map_dbl(semantic_coherence, mean)) %>%
  gather(Metric, Value, -K) %>%
  ggplot(aes(K, Value, color = Metric)) +
  geom_line(linewidth = 1.5, alpha = 0.7, show.legend = FALSE) +
  facet_wrap(~Metric, scales = "free_y") +
  labs(x = "K (number of topics)",
       y = NULL,
       title = "Model diagnostics by number of topics")
```

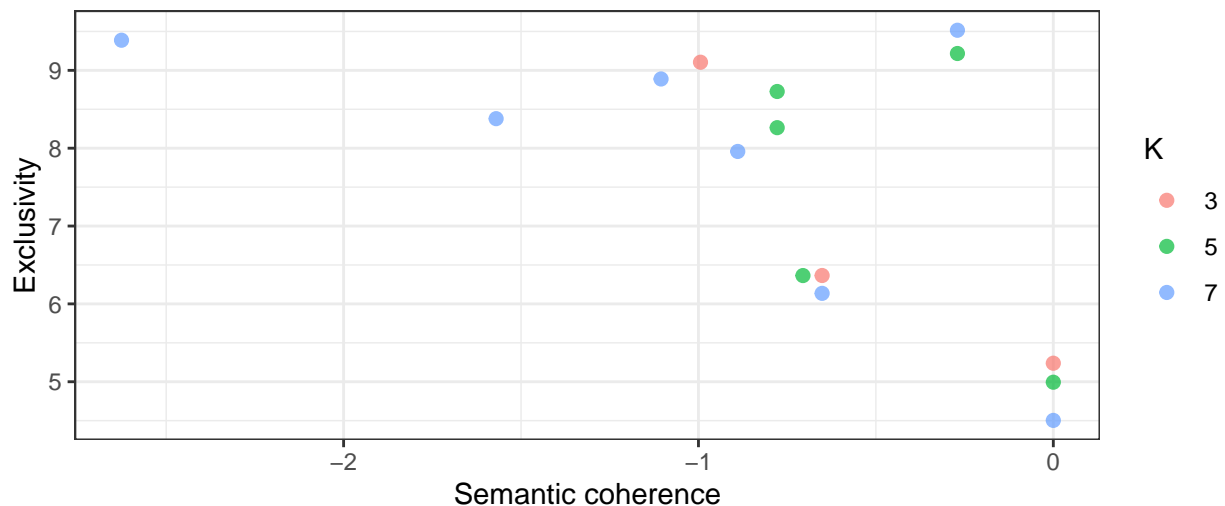
## Model diagnostics by number of topics



Semantic coherence is maximized when the most probable words in a given topic frequently co-occur together, and it's a metric that correlates well with human judgment of topic quality. Having high semantic coherence is relatively easy, though, if you only have a few topics dominated by very common words, so you want to look at both semantic coherence and exclusivity of words to topics. It's a tradeoff. Read more about semantic coherence in the original paper about it (<https://dl.acm.org/citation.cfm?id=2145462>).

```
df_diagnostics %>%
  select(K, exclusivity, semantic_coherence) %>%
  unnest(cols = c(exclusivity, semantic_coherence)) %>%
  filter(K %in% c(3, 5, 7)) %>%
  mutate(K = as.factor(K)) %>%
  ggplot(aes(semantic_coherence, exclusivity, color = K)) +
  geom_point(size = 2, alpha = 0.7) +
  labs(x = "Semantic coherence",
       y = "Exclusivity",
       title = "Comparing exclusivity and semantic coherence")
```

## Comparing exclusivity and semantic coherence



In choosing our final topic model, we should go with a K value that is statistically sound, but also contextually (in this case, biologically) meaningful and informative

## Finalize model

```
topic_model <- stm(sparse_mat,
                  K = 6, # K is the number of 'topics' (clusters)
                  verbose = FALSE)

p_beta <- tidy(topic_model, matrix = "frex") %>%
  left_join(tidy(topic_model)) %>%
  group_by(topic) %>%
  slice_head(n = 20) %>%
  mutate(topic = paste0("topic ", topic)) %>%

  ggplot(aes(x = beta, y = reorder(term, beta), fill = topic)) +
  geom_col() +
  labs(y = "Word") +
  facet_wrap(vars(topic), scales = "free", nrow = 1)

# Topic probabilities
group_gamma <- tidy(
  topic_model,
  matrix = "gamma",
  document_names = rownames(sparse_mat)) %>%
  mutate(chapter = factor(document)) %>%
  dplyr::select(chapter, topic, gamma) %>%
  mutate(topic = factor(topic))
)

p_gamma <- group_gamma %>%
  filter(gamma > 0.01) %>%
  ggplot(aes(x = topic, y = gamma, color = topic)) +
  geom_point(aes(alpha = gamma)) +
  geom_text_repel(aes(alpha = gamma, label = chapter), size = 3,
                 max.overlaps = 50) +

  labs(y = expression(gamma), x = "") +
  ggtitle("Corpus (character) strength of association with each topic") +
  theme(legend.position = "none",
        axis.text.x = element_blank())

(p_gamma / p_beta)
```

Corpus (character) strength of association with each topic

